

Model matters with restricted Boltzmann machines

Andee Kaplan

Colorado State University
andee.kaplan@colostate.edu

November 20, 2019

Slides available at <http://bit.ly/kaplan-silo>

Joint work with D. Nordman and S. Vardeman

What is this?

A restricted Boltzman machine (RBM) is an undirected probabilistic graphical model with

- 1 two layers of random variables - one hidden and one visible
- 2 conditional independence within a layer (Smolensky 1986)




Figure 1: Hidden nodes are indicated by white circles and the visible nodes are indicated by blue circles.

How is it used?

- Supervised learning, specifically image classification




Figure 2: Image classification using a RBM: each image pixel comprises a node in the visible layer, \mathcal{V} and the output of the RBM is used to create features passed to a supervised learning algorithm.

Joint distribution

- $\mathbf{x} = (h_1, \dots, h_{n_H}, v_1, \dots, v_{n_V})$ represents visible and hidden nodes in a RBM
- Each single “binary” random variable, visible v_i or hidden h_j , takes values in a common coding set
 - $\mathcal{C} = \{0, 1\}$ or $\mathcal{C} = \{-1, 1\}$.
- A parametric form for probabilities

$$f_{\theta}(\mathbf{x}) = \frac{\exp \left(\sum_{i=1}^{n_V} \sum_{j=1}^{n_H} \theta_{ij} v_i h_j + \sum_{i=1}^{n_V} \theta_{v_i} v_i + \sum_{j=1}^{n_H} \theta_{h_j} h_j \right)}{\gamma(\theta)}$$

where

$$\gamma(\theta) = \sum_{\mathbf{x} \in \mathcal{C}^{n_H+n_V}} \exp \left(\sum_{i=1}^{n_V} \sum_{j=1}^{n_H} \theta_{ij} v_i h_j + \sum_{i=1}^{n_V} \theta_{v_i} v_i + \sum_{j=1}^{n_H} \theta_{h_j} h_j \right)$$

Deep learning

- Stacking layers of RBMs in a deep architecture
- Proponents claim the ability to learn "internal representations that become increasingly complex, which is considered to be a promising way of solving object and speech recognition problems" (Salakhutdinov and Hinton 2009, pp. 450).




Figure 3: Three layer deep Boltzmann machine, with visible-to-hidden and hidden-to-hidden connections but no within-layer connections.

Why do I care?

- ① The model properties are largely unexplored in the literature
- ② The commonly cited fitting methodology remains heuristic-based (Hinton, Osindero, and Teh 2006)

We want to

- ① Provide steps toward understanding properties of the model class from the perspective of statistical theory
- ② Explore the possibility of a rigorous fitting methodology

Properties of RBM Model Class

Degeneracy, instability, and uninterpretability. Oh my!

The highly flexible nature of a RBM ($n_H + n_V + n_H * n_V$ parameters) makes at least three kinds of potential model impropriety of concern

- 1 *degeneracy*
- 2 *instability*, and
- 3 *uninterpretability*

A model should “provide an explanation of the mechanism underlying the observed phenomena” (G. E. P. Box 1967).

RBM's often

- fail to generate data with realistic variability and thus an unsatisfactory conceptualization of the data generation process (Li 2014)
- exhibit model instability (over-sensitivity) (Szegedy et al. 2013; Nguyen, Yosinski, and Clune 2014)

Near-degeneracy

Definition (Model Degeneracy)

A disproportionate amount of probability is placed on only a few elements of the sample space, $\mathcal{C}^{n_H+n_V}$, by the model.

RBM models exhibit *near-degeneracy* when random variables in

$$Q_{\theta}(\mathbf{x}) = \sum_{i=1}^{n_V} \sum_{j=1}^{n_H} \theta_{ij} v_i h_j + \sum_{i=1}^{n_V} \theta_{v_i} v_i + \sum_{j=1}^{n_H} \theta_{h_j} h_j,$$

have a mean vector $\boldsymbol{\mu}(\boldsymbol{\theta})$ close to the boundary of the convex hull of $\mathcal{T} = \{\mathbf{t}(\mathbf{x}) : \mathbf{x} \in \mathcal{C}^{n_H+n_V}\}$ (Handcock 2003), where

$$\mathbf{t}(\mathbf{x}) = \{v_1, \dots, v_{n_V}, h_1, \dots, h_{n_H}, v_1 h_1, \dots, v_{n_V} h_{n_H}\}$$

and

$$\boldsymbol{\mu}(\boldsymbol{\theta}) = \mathbb{E}_{\theta} \mathbf{t}(\mathbf{X})$$

Instability

Definition (Instability)

Characterized by excessive sensitivity in the model, where small changes in the components of data outcomes, \mathbf{x} , lead to substantial changes in probability.

- Concept of model deficiency related to *instability* for a class of exponential families of distributions (Schweinberger 2011)
- For the RBM, consider how model incorporates more visibles
 - Model parameters in a longer sequence $\theta_{n_V} \in \mathbb{R}^{n_V + n_H + n_V * n_H}$, $n_V \geq 1$
 - May also arbitrarily expand the number of hidden variables used

Unstable RBMs

Definition (S-unstable RBM)

A RBM model formulation is *S-unstable* if

$$\lim_{n_V \rightarrow \infty} \frac{1}{n_V} \text{LREP}(\theta_{n_V}) = \infty.$$

where

$$\text{LREP}(\theta_{n_V}) = \log \left[\frac{\max_{(v_1, \dots, v_{n_V}) \in \mathcal{C}^{n_V}} P_{\theta_{n_V}}(v_1, \dots, v_{n_V})}{\min_{(v_1, \dots, v_{n_V}) \in \mathcal{C}^{n_V}} P_{\theta_{n_V}}(v_1, \dots, v_{n_V})} \right] \quad (1)$$

S-unstable RBM models are undesirable for several reasons - small changes in data outcomes can lead to overly-sensitive changes in probability.

One-pixel change

Consider the biggest log-probability ratio for a one-pixel (one component) change in visibles (data outcomes)

$$\Delta(\theta_{n_V}) \equiv \max \left\{ \log \frac{P_{\theta_{n_V}}(v_1, \dots, v_{n_V})}{P_{\theta_{n_V}}(v_1^*, \dots, v_{n_V}^*)} \right\},$$

where (v_1, \dots, v_{n_V}) & $(v_1^*, \dots, v_{n_V}^*) \in \mathcal{C}^{n_V}$ differ by exactly one component

Result

Let $c > 0$ and fix an integer $n_V \geq 1$. If $\frac{1}{n_V} LREP(\theta_{n_V}) > c$, then $\Delta(\theta_{n_V}) > c$.

If the $n_V^{-1} LREP(\theta_{n_V})$ is too large, then a RBM model will exhibit large probability shifts for very small changes in the data configuration.

Tie to degeneracy

Define an arbitrary modal set of possible outcomes (i.e. set of highest probability outcomes) for a given $0 < \epsilon < 1$ as

$$M_{\epsilon, \theta_{n_V}} \equiv \left\{ \mathbf{v} \in \mathcal{C}^{n_V} : \log P_{\theta_{n_V}}(\mathbf{v}) > (1 - \epsilon) \max_{\mathbf{v}^*} P_{\theta_{n_V}}(\mathbf{v}^*) + \epsilon \min_{\mathbf{v}^*} P_{\theta_{n_V}}(\mathbf{v}^*) \right\}$$

Result

For an S -unstable RBM model, and for any given $0 < \epsilon < 1$, $P_{\theta_{n_V}} \left((v_1, \dots, v_{n_V}) \in M_{\epsilon, \theta_{n_V}} \right) \rightarrow 1$ holds as $n_V \rightarrow \infty$.

- All probability will stack up on mode sets or potentially those few outcomes with the highest probability

Uninterpretability

Definition (Uninterpretability)

Characterized by marginal mean-structure (controlled by main effect parameters $\theta_{v_i}, \theta_{h_j}$) not being maintained in the model due to dependence (interaction parameters θ_{ij}) (Kaiser 2007).

- Model expectations, $E[\mathbf{X}|\boldsymbol{\theta}]$
- Expectations given independence, $E[\mathbf{X}|\boldsymbol{\theta}^*]$, where $\boldsymbol{\theta}^*$ matches $\boldsymbol{\theta}$ for all main effects but otherwise has $\theta_{ij} = 0$ for $i = 1, \dots, n_V, j = 1, \dots, n_H$
- If $|E[\mathbf{X}|\boldsymbol{\theta}] - E[\mathbf{X}|\boldsymbol{\theta}^*]|$ is large then the RBM with parameter vector $\boldsymbol{\theta}$ is *uninterpretable*

FOES models

Statements about near-degeneracy, instability, and uninterpretability hold for a general class of models that includes the RBM.

- $\mathbf{X}_N = (X_1, \dots, X_N)$ collection of discrete random variables with a finite sample space
- \mathcal{X}^N finite sample space, represented as some N -fold Cartesian product
- $P_{\theta_{q(N)}}$ probability model on \mathcal{X}^N for each N
- Assume that the model support of $P_{\theta_{q(N)}}$ is the sample space \mathcal{X}^N .

Definition

This framework produces *Finite Outcome Everywhere Supported (FOES)* models, $P_{\theta_{q(N)}}$, indexed by a defining sequence of parameters $\theta_{q(N)}$, to describe data \mathbf{X}_N of any given sample size $N \geq 1$.

(Kaplan, Nordman, and Vardeman 2019a)

Examples

- Discrete exponential family models
- Discrete restricted Boltzmann machines
- Binary deep learning models
 - ① **Deep Boltzmann machine (DBM)**. Conditional independence within all layers, stacked RBM models with conditional dependence between neighboring layers
 - ② **Deep belief network (DBN)**. Multiple layers of latent random variables stacked in a deep architecture with no conditional dependence between layers; all but the last stacked layer are Bayesian networks (directed dependence), rather than RBMs

Visualizing properties in manageable (a.k.a. small) RBMs

- To explore the effects of RBM parameters θ on *near-degeneracy*, *instability*, and *uninterpretability*, consider models of small size
- For $n_H, n_V \in \{1, \dots, 4\}$, sample 100 values of θ
 - ① Split θ into $\theta_{interaction}$ and θ_{main}
 - ② Allow the two types of terms to have varying average magnitudes, $||\theta_{main}||/(n_H + n_V)$ and $||\theta_{interaction}||/(n_H * n_V)$
 - ③ Average magnitudes vary on a grid between 0.001 and 3 with 24 breaks, yielding 576 grid points
- Calculate metrics of model impropriety, $\mu(\theta)$, $LREP(\theta)/n_V$, and the coordinates of $|E[X|\theta] - E[X|\theta^*]|$.
- In the case of *near-degeneracy*, classify each model as near-degenerate or “viable” based on the distance of $\mu(\theta)$ from the boundary of the convex hull of \mathcal{T}

Simulation results

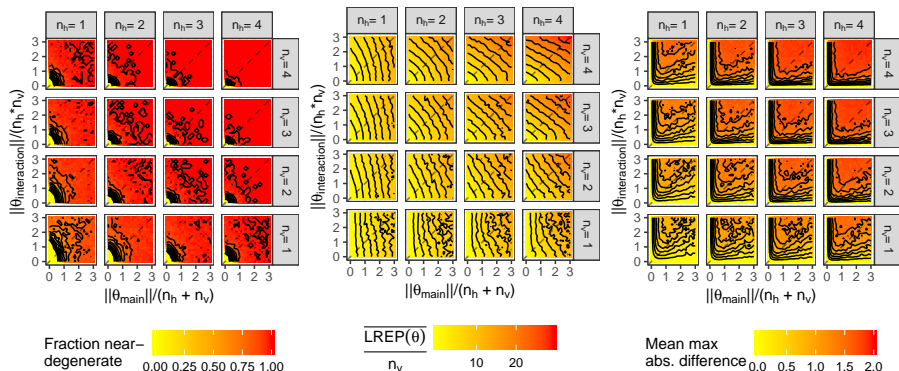


Figure 4: The fraction of models that were near-degenerate (left), the sample mean value of $\text{LREP}(\theta)/n_v$ (middle), and the sample mean of the maximum component of the absolute difference between the model expectation vector, $E[\mathbf{X}|\theta]$, and the expectation vector given independence, $E[\mathbf{X}|\theta^*]$ (right).

Rigorous Fitting Methodology

Model fitting

- ① Computational concerns: Fitting a RBM via maximum likelihood (ML) methods infeasible due to the intractability of the normalizing term $\gamma(\theta)$
 - Ad hoc methods used to avoid this problem with stochastic ML
 - Employ a small number of MCMC draws to approximate $\gamma(\theta)$
- ② Model parameterization concerns: With enough hidden units,
 - Potential to re-create any distribution for the data (Le Roux and Bengio 2008; Montufar and Ay 2011; and Montúfar, Rauh, and Ay 2011)
 - The model for the cell probabilities that has the highest likelihood over *all possible model classes* is the empirical distribution
 - The RBM model ensures that this empirical distribution can be arbitrarily well approximated
 - When empirical distribution contains empty cells, ML will chase parameters to ∞ in order to zero out corresponding RBM cell probabilities

Bayesian methods

- Consider what might be done in a principled manner, small test
- To avoid model impropriety, avoid parts of the parameter space $\mathbb{R}^{n_V + n_H + n_V * n_H}$ leading to *near-degeneracy*, *instability*, and *uninterpretability*.
 - Shrink θ toward $\mathbf{0}$
 - 1 Specify priors that place low probability on large values of $||\theta||$
 - 2 Shrink $\theta_{interaction}$ more than θ_{main}
- Consider a test case with $n_V = n_H = 4$
 - θ chosen as a sampled value from a grid point in figure 4 with $< 5\%$ near-degeneracy (not near the convex hull of the sufficient statistics)
 - simulate $n = 5,000$ as a training set and fit the RBM using three Bayes methodologies

Fitting methodologies

① A “trick” prior (BwTPLV)

- Cancel out normalizing term in the likelihood
- Resulting full conditionals of θ are multivariate Normal
- h_j are carried along as latent variables

$$\pi(\theta) \propto \gamma(\theta)^n \exp \left(-\frac{1}{2C_1} \theta'_{main} \theta_{main} - \frac{1}{2C_2} \theta'_{interaction} \theta_{interaction} \right),$$

where $C_2 < C_1$ (Li 2014)

Fitting methodologies (cont'd)

② *A truncated Normal prior (BwTNLV)*

- Independent spherical normal distributions as priors for θ_{main} and $\theta_{interaction}$
 - $\sigma_{interaction} < \sigma_{main}$
 - *truncated* at $3\sigma_{main}$ and $3\sigma_{interaction}$, respectively
- Simulation from the posterior using a geometric adaptive MH step (Zhou 2014)
- h_j are carried along in the MCMC implementation as latent variables

③ *A truncated Normal prior and marginalized likelihood (BwTNML)*

- Marginalize out \mathbf{h} in $f_{\theta}(\mathbf{x})$
- Use the truncated Normal priors applied to the marginal probabilities for visible variables (recall visibles are the observed data, hiddens are not)

Hyperparameters

Table 1: The values used for the hyperparameters for all three fitting methods. A rule of thumb is imposed which decreases prior variances for the model parameters as the size of the model increases and also shrinks $\theta_{interaction}$ more than θ_{main} . The common C defining C_1 and C_2 in the BwTPLV method is chosen by tuning.

Method	Hyperparameter	Value
BwTPLV	C_1	$\frac{C}{n} \frac{1}{n_H + n_V}$
	C_2	$\frac{C}{n} \frac{1}{n_H * n_V}$
BwTNLV	σ_{main}^2	$\frac{1}{n_H + n_V}$
	$\sigma_{interaction}^2$	$\frac{1}{n_H * n_V}$
BwTNML	σ_{main}^2	$\frac{1}{n_H + n_V}$
	$\sigma_{interaction}^2$	$\frac{1}{n_H * n_V}$

Mixing

The BwTNLV (2) and the BwTNML method (3) are drawing from the same stationary posterior distribution for images.




Figure 5: The autocorrelation functions (ACF) for the posterior probabilities of all $2^4 = 16$ possible outcomes for the vector of 4 visibles assessed at multiple lags for each method with BwTNLV in black and BwTNML in red.

Effective sample size

- Overlapping blockmeans approach (Gelman, Shirley, and others 2011)
 - Crude estimate for the asymptotic variance of the probability of each image
 - Compare it to an estimate of the asymptotic variance assuming IID draws from the target distribution

Table 2: The effective sample sizes for a chain of length $M = 1000$ regarding all 16 probabilities for possible vector outcomes of visibles. BwTNLV would require at least 4.7 times as many MCMC iterations to achieve the same amount of effective information about the posterior distribution.

Outcome	BwTNLV	BwTNML	Outcome	BwTNLV	BwTNML
1	73.00	509.43	9	83.47	394.90
2	65.05	472.51	10	95.39	327.35
3	87.10	1229.39	11	70.74	356.56
4	72.64	577.73	12	81.40	338.30
5	71.67	452.01	13	105.98	373.59
6	66.49	389.78	14	132.61	306.91
7	84.30	660.37	15	82.15	365.30
8	75.46	515.09	16	98.05	304.57

Posterior distributions of images

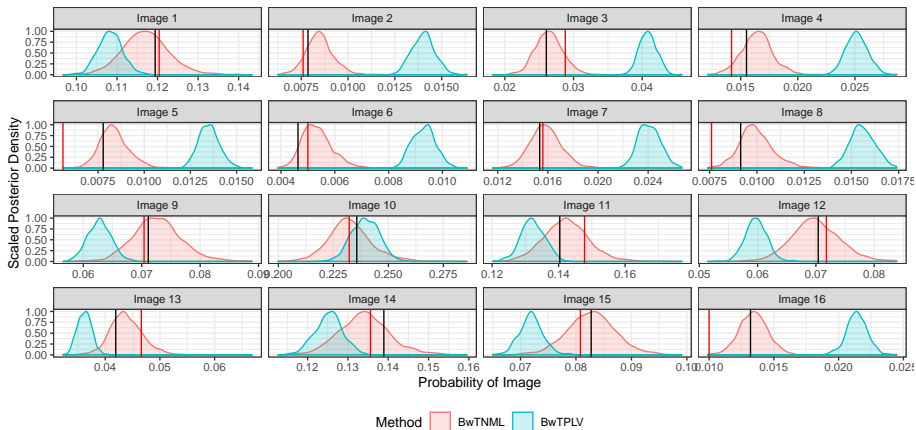


Figure 6: Posterior probabilities of $16 = 2^4$ possible realizations of 4 visibles using two of the three Bayesian fitting techniques, BwTPLV and BwTNML. Black lines show true probabilities of each vector of visibles based on the parameters used to generate the training data while red lines show the empirical distribution.

Wrapping up

- RBMs shown success for classification, but concerning as statistical models due to *near-degeneracy*, *S-instability*, and *uninterpretability*
- Rigorous fitting methodology is difficult
 - Numerical complications in likelihood maximization may occur due to degeneracy and S-instability when all probability is pushed to opposite extremes in the sample space
 - For Bayesian inference, S-instability may hinder effective chain mixing
- RBMs questionably useful as any distribution for the visibles can be approximated arbitrarily well
 - The empirical distribution of visibles is the best fitting model for observed cell data
 - There can be no “smoothed distribution” achieved in a RBM model of sufficient size with a rigorous likelihood-based method

Skeptical that any model built using RBMs (e.g., DBM) can achieve useful **inference** in principled way without limiting flexibility of the fitted model

Relevant papers

- 1 Kaplan, Andee, Daniel Nordman, and Stephen Vardeman. “Properties and Bayesian fitting of restricted Boltzmann machines.” *Statistical Analysis and Data Mining: The ASA Data Science Journal* 12.1 (2019): 23-38.
- 2 Kaplan, Andee, Daniel J. Nordman, and Stephen B. Vardeman. “On the S-instability and degeneracy of discrete deep learning models.” *Information and Inference: A Journal of the IMA* (2019).

Thank you

- Slides – <http://bit.ly/kaplan-silo>
- Contact
 - Email – andee.kaplan@colostate.edu
 - Twitter – <http://twitter.com/andee Kaplan>
 - GitHub – <http://github.com/andee k>

Appendix: Parameters used

Table 3: Parameters used to fit a test case with $n_v = n_h = 4$. This parameter vector was chosen as a sampled value of θ that was not near the convex hull of the sufficient statistics for a grid point in figure 4 with $< 5\%$ near-degeneracy.

Parameter	Value	Parameter	Value	Parameter	Value
θ_{v1}	-1.1043760	θ_{11}	-0.0006334	θ_{31}	-0.0038301
θ_{v2}	-0.2630044	θ_{12}	-0.0021401	θ_{32}	0.0032237
θ_{v3}	0.3411915	θ_{13}	0.0047799	θ_{33}	0.0020681
θ_{v4}	-0.2583769	θ_{14}	0.0025282	θ_{34}	0.0041429
θ_{h1}	-0.1939302	θ_{21}	0.0012975	θ_{41}	0.0089533
θ_{h2}	-0.0572858	θ_{22}	0.0000253	θ_{42}	-0.0042403
θ_{h3}	-0.2101802	θ_{23}	-0.0004352	θ_{43}	-0.0000480
θ_{h4}	0.2402456	θ_{24}	-0.0086621	θ_{44}	0.0004767

References I

- Gelman, Andrew, Kenneth Shirley, and others. 2011. "Inference from Simulations and Monitoring Convergence." *Handbook of Markov Chain Monte Carlo*, 163–74.
- G. E. P. Box, W. J. Hill. 1967. "Discrimination Among Mechanistic Models." *Technometrics* 9 (1). [Taylor & Francis, Ltd., American Statistical Association, American Society for Quality]: 57–71.
- Handcock, Mark S. 2003. "Assessing Degeneracy in Statistical Models of Social Networks." Center for Statistics; the Social Sciences, University of Washington. <http://www.csss.washington.edu/>.
- Hinton, Geoffrey E, Simon Osindero, and Yee-Whye Teh. 2006. "A Fast Learning Algorithm for Deep Belief Nets." *Neural Computation* 18 (7). MIT Press: 1527–54.
- Kaiser, Mark S. 2007. "Statistical Dependence in Markov Random Field Models." *Statistics Preprints* Paper 57. Digital Repository @ Iowa State University. http://lib.dr.iastate.edu/stat_las_preprints/57/.
- Kaplan, Andee, Daniel J Nordman, and Stephen B Vardeman. 2019a. "On the S-instability and degeneracy of discrete deep learning models." *Information and Inference: A Journal of the IMA*, November. <https://doi.org/10.1093/imaiai/iaz022>.
- Kaplan, Andee, Daniel Nordman, and Stephen Vardeman. 2019b. "Properties and Bayesian Fitting of Restricted Boltzmann Machines." *Statistical Analysis and Data Mining: The ASA Data Science Journal* 12 (1): 23–38. <https://doi.org/10.1002/sam.11396>.
- Le Roux, Nicolas, and Yoshua Bengio. 2008. "Representational Power of Restricted Boltzmann Machines and Deep Belief Networks." *Neural Computation* 20 (6). MIT Press: 1631–49.
- Li, Jing. 2014. "Biclustering Methods and a Bayesian Approach to Fitting Boltzmann Machines in Statistical Learning." PhD thesis, Iowa State University; Graduate Theses; Dissertations. <http://lib.dr.iastate.edu/etd/14173/>.
- Montufar, Guido, and Nihat Ay. 2011. "Refinements of Universal Approximation Results for Deep Belief Networks and Restricted Boltzmann Machines." *Neural Computation* 23 (5). MIT Press: 1306–19.

References II

- Montúfar, Guido F, Johannes Rauh, and Nihat Ay. 2011. "Expressive Power and Approximation Errors of Restricted Boltzmann Machines." In *Advances in Neural Information Processing Systems*, 415–23. NIPS.
- Nguyen, Anh Mai, Jason Yosinski, and Jeff Clune. 2014. "Deep Neural Networks Are Easily Fooled: High Confidence Predictions for Unrecognizable Images." *arXiv Preprint arXiv:1412.1897*. <http://arxiv.org/abs/1412.1897>.
- Salakhutdinov, Ruslan, and Geoffrey E Hinton. 2009. "Deep Boltzmann Machines." In *International Conference on Artificial Intelligence and Statistics*, 448–55. AI & Statistics.
- Schweinberger, Michael. 2011. "Instability, Sensitivity, and Degeneracy of Discrete Exponential Families." *Journal of the American Statistical Association* 106 (496). Taylor & Francis: 1361–70.
- Smolensky, Paul. 1986. "Information Processing in Dynamical Systems: Foundations of Harmony Theory." DTIC Document.
- Szegedy, Christian, Wojciech Zaremba, Ilya Sutskever, Joan Bruna, Dumitru Erhan, Ian J. Goodfellow, and Rob Fergus. 2013. "Intriguing Properties of Neural Networks." *arXiv Preprint arXiv:1312.6199*. <http://arxiv.org/abs/1312.6199>.
- Zhou, Wen. 2014. "Some Bayesian and Multivariate Analysis Methods in Statistical Machine Learning and Applications." PhD thesis, Iowa State University; Graduate Theses; Dissertations. <http://lib.dr.iastate.edu/etd/13816/>.